

Rohith Kota

• San Francisco, CA • Kotarohit14@gmail.com • +1 (857)335-3738 • [LinkedIn](#) • Open to Relocate

Data & AI Engineer | LLM Applications | Scalable ETL | Lakehouse Architect

“Data & AI Engineer with 5+ years of experience managing **petabyte-scale data**, building **scalable LLM-powered applications**, **Data Lakehouses**, and over **200+** ETL and ML pipelines to enhance data reliability and accessibility.”

SKILLS

Programming Languages: Python (PySpark, Pandas, NumPy), SQL (T-SQL, PL/SQL, SparkSQL), Scala

Big Data Technologies: Spark, Airflow, Kafka, HDFS, Hive, Presto, Hadoop

ETL & Visualization: Airflow, Azure Data Factory, Dagster, DBT, Databricks, Fivetran, Power BI, Tableau, Looker

Databases & Storage: Snowflake, PostgreSQL, Redshift, ADLS, Azure Synapse, S3, MySQL, MSSQL, MongoDB (NoSQL)

Gen AI: Prompt Engineering, Finetuning, Langchain, HuggingFace, CrewAI, RAG, LangGraph, LLM Agents

Other Tools: Git, Agile, JIRA, Docker, VS code, Kubernetes, CI/CD Pipelines, MLflow

PROJECTS

Multi AI Agent System- Large language Model (ChatGPT – 4)

- Designed scalable AI agent architecture with LangChain and OpenAI models to answer complex user queries.
- Built PythonExecAgent and DateTimeAgent for contextual execution; integrated GPT-3 and BERT for comprehension.
- Coordinated all agents via a Central Orchestration Agent (COA), improving task resolution accuracy by 30%.

Knowledge Graph for RAG

- Streamlined RAG pipeline with OpenAI embeddings and LangChain stuff chain to generate accurate responses.
- Integrated Cypher query generation and Knowledge Graph traversal using GraphCypherQA chain.
- Enabled non-technical users to interact via Gradio UI; improved retrieval precision and response relevance.

WORK EXPERIENCE

Data Engineer

Aug 2024 – Present

META

Menlo Park, CA

- Engineered end-to-end Ads metrics data foundations for 0→1 product launches; unlocked \$300M in incremental ad revenue.
- Automated growth analytics via Airflow, Hive, and standardized ETL, Adopted by **9+** product teams reducing dev time by 80%.
- Created attribution parameter matching system in Presto for 250B+ records; standardized **integration** of [Google Analytics](#) and **Adobe** logs via **API**. Uplifted customer’s conversion rate by **5%** and **12%** increase in their **marketing** spend.
- Orchestrated **Spark + Hive**-based data pipeline processing **200 TB/day** to support **ML training** and **A/B testing** to subsidize **high-CPA** ad campaigns, leading to a **\$150M** boost in advertiser spend and a **20%** increase in conversions.

Data & AI Engineer

Jun 2023 – Jul 2024

T-Mobile

Seattle, WA

- Deployed **LLM-powered** analytics assistant using **LangChain + Azure OpenAI**; integrated **semantic SQL retrieval** across **Presto** and **Snowflake**, accelerating ad hoc business insights by 40% and enabling **non-SQL** users to **self-serve**.
- Built **prompt monitoring** infrastructure using **Prometheus** and **Airflow**; automated **re-embedding** using **FAISS**.
- Implemented a **Medallion-layer** on **Azure (ADF + Databricks over ADLS)**, orchestrating **batch** pipelines that feed governed Gold layers to **Power BI** dashboards—cut **ETL latency 20%** and time-to-insight **30%**.
- Identified vendor charge discrepancies in **Snowflake** using **DBT**; proposed modifications saved **16%** in costs.

Data & AI Research Asst.

Nov 2022 – Jun 2023

Drexel University

Philadelphia, PA

- Applied Retrieval-Augmented Generation using LangChain + OpenAI for academic search assistant.
- Implemented GraphCypherQA chain for complex research queries, enabling deeper insights.
- Rolled out chatbot with Gradio interface; supported **50+** graduate researchers in complex query resolution.

Analytical Engineer

Jun 2020 - Aug 2022

HighRadius

Hyderabad, TG

- Developed **AI AR** automation product generating **\$300K/quarter**; utilized by **80+** Businesses and reduced DSO by **30%**.
- Integrated deduction validity **ML** models and automated workflows, improving recovery rates by **30%**.
- Deployed **Lambda + API** Gateway to monitor credit risk; reduced bad debt 10%, onboarding time by **70%**.
- Fabricated a **Data Warehouse** for an E-fitness client with **100+ tables** and composed **25+ ETL** Pipelines using **Azure Data Factory** from **MongoDB, MySQL**, App Store Connect **API** to **MSSQL** Server and **Tableau** dashboards; cut data redundancy by 38%.
- Designed scalable, fault-tolerant data pipelines using **Kafka** and Amazon **Kinesis** for **real-time** data ingestion into **Redshift**, achieving a **25%** reduction in latency through parallel processing techniques.

EDUCATION

Drexel University, Le Bow College of Business (**3.89/4.0**)

Master of Science, Business Analytics

Philadelphia, PA